

Mauve Assembly Metrics

Aaron Darling^{1*}, Andrew Tritt¹, Jonathan A. Eisen^{1,2,3} and Marc T. Facciotti^{1,4}

¹Genome Center, ² Dept. of Evolution and Ecology, ³ Medical Microbiology and Immunology, ⁴ Biomedical Engineering, University of California-Davis, Davis, CA 95616.

Associate Editor: Dr. Alex Bateman

ABSTRACT

Summary: High throughput DNA sequencing technologies have spurred the development of numerous novel methods for genome assembly. With few exceptions, these algorithms are heuristic and require one or more parameters to be manually set by the user. One approach to parameter tuning involves assembling data from an organism with an available high quality reference genome, and measuring assembly accuracy using some metrics.

We developed a system to measure assembly quality under several scoring metrics, and to compare assembly quality across a variety of assemblers, sequence data types, and parameter choices. When used in conjunction with training data such as a high quality reference genome and sequence reads from the same organism, our program can be used to manually identify an optimal sequencing and assembly strategy for *de novo* sequencing of related organisms.

Availability: GPL source code and a usage tutorial is at <http://ngopt.googlecode.com>

Contact: aaron.darling@ucdavis.edu

1 INTRODUCTION

Given high throughput sequencing data, most current genome assemblers apply deterministic heuristics to infer the genome sequence. Usually a variety of parameters can be used to control the heuristic, for which the optimal combination of values may not be obvious. Given a training dataset consisting of high quality reference genomes and sequence reads generated from those genomes, it may be possible to manually or automatically select a good set of assembly parameters. A key requirement for this task is a means to quantify the accuracy of an assembly.

Measuring the accuracy with which an assembly reconstructs the reference genome presents another inference problem. It is usually unknown which part of the inferred assembly corresponds to which part of the reference genome. We must somehow map parts of the assembly back onto the reference genome through sequence alignment, which usually takes one of two forms: local alignment, exemplified by algorithms like BLAST (Altschul *et al.*, 1997), and whole genome alignment with algorithms like MUMmer (Kurtz *et al.*, 2004) or Mauve (Darling *et al.*, 2004).

We introduce a new set of assembly accuracy metrics based on the *progressiveMauve* genome aligner (Darling *et al.*, 2010). In our method, the assembly contigs and/or scaffolds are first reordered

to match a reference genome with the Mauve Contig Mover (Rissman *et al.*, 2009). The ordered, aligned assembly is then compared to the reference to identify differences.

Our method is most closely related to MUMmer's *dnadiff* program, which can measure assembly errors using genome alignment (Phillippy *et al.*, 2008). We, however, use a different alignment heuristic and evaluate some new types of error such as rearrangement distance. Several ongoing efforts are directed at measuring assembly accuracy on particular datasets, including the Assemblathon, GAGE, and dnGASP. These initiatives use tools like *dnadiff*, the Mauve Assembly Metrics, and others.

2 METHODS

In the present work we summarize differences in a pairwise alignment of the assembly and reference genome (e.g. as computed by *progressiveMauve* (Darling *et al.*, 2010)). We illustrate this process by way of example. Given the following reference genome and assembled genome:

Reference: AGGCTAGCGCGCGATTAGGATC

Assembly: AGTAGCGGGCCGATTAGGATC

A genome alignment of the reference and assembly might look like:

Reference: AGGCTAGCGCG-CGATTAGGATC

Assembly: AG--TAGCGGGCCGATTAGGATC

From this alignment, we would calculate the assembly scoring metrics as follows (not an exhaustive list of metrics):

- Miscalled bases: 2 (C→G and G→A)
- Uncalled bases: 1 (N)
- Extra bases: 1 (Insertion of C in assembly)
- Missing bases: 2 (Deletion of GC in assembly)
- Number of extra segments: 1
- Number of missing segments: 1

In addition to metrics summarizing the number of base miscalls, missing, and extra segments (each also evaluated by *dnadiff*), our method produces a variety of other metrics. The location of miscalled bases, missing segments, and extra segments is exported to a tab-delimited text file for subsequent analysis. GC content of the missing and extra regions is also exported. Misassemblies are identified as rearrangement breakpoints inside of contigs. The double cut and join (DCJ) distance (Bergeron *et al.*, 2006) between the assembly and reference is calculated to estimate the combined effect of misassembly and lack-of-assembly errors (excess contig breaks) on

*to whom correspondence should be addressed

Table 1. Mauve assembly metrics for three assemblies of *Haloferax volcanii* DS2.

Metric	volc454	volcV	volcIDBA
Scaffold count	157	1394	50
Miscalled bases	81	948	235
Uncalled bases	0	53899	15188
Extra bases	0.04%	10.8%	2.54%
Missing bases	3.13%	5.87%	2.71%
Extra segments	43	1079	262
Missing segments	117	1144	192
DCJ Distance	114	909	61
Intact CDS	99.3%	87.8%	97.3%

rearrangement distance. Finally each protein coding sequence in the reference genome is checked in the assembly for whether it yields an intact coding sequence, with types and location of substitution and frameshift errors reported.

2.1 Assembling genomes of Haloarchaea

In an ongoing effort, we are sequencing *de novo* the genomes of 60 halophilic archaea. Four of these organisms have high quality reference genomes completed independently of our project. We elected to demonstrate our new assembly metrics on one of these organisms, *Haloferax volcanii* strain DS2. This organism has a 4.0 Mbp genome organized into five circular replicons with about 100 repetitive IS elements of 1-2 Kbp each (Hartman *et al.*, 2010). Using 454 and Illumina resequencing data, we generated three different assemblies to compare with our software. The assemblies are named **volc454**, **volcV**, and **volcIDBA** (see Supplementary Material for sequencing and assembly details).

We scored each assembly against the reference genome using the aforementioned method. An overview of each assembly's metrics is given in Table 1. The location of assembly errors is mapped on the *Haloferax volcanii* DS2 reference genome in Figure 1. Finally, Figure 2 illustrates that each sequencing and assembly strategy appears to have bias in the direction of erroneous base calls.

3 DISCUSSION

The assembly metrics we describe illustrate substantial differences between sequencing and assembly strategies. For example, the volc454 assembly captured nearly all coding genes in the reference genome, but had a high scaffold count relative to the volcIDBA assembly. Striking an ideal balance between assembly error types, rates and sequencing cost is an exercise left for users of our software.

When a finished reference genome is available and has been resequenced, the assembly metrics calculated by our system can be used to guide selection of sequencing strategy and tune assembly parameters. The reported metrics may form the basis for a future automated system to perform supervised machine learning of assembly parameters by conducting a parameter sweep over a large number of assembly strategies.

Finally, we note that genome alignment algorithms are not perfect and some differences between the assembly and the reference may be due to alignment error and not true assembly errors.

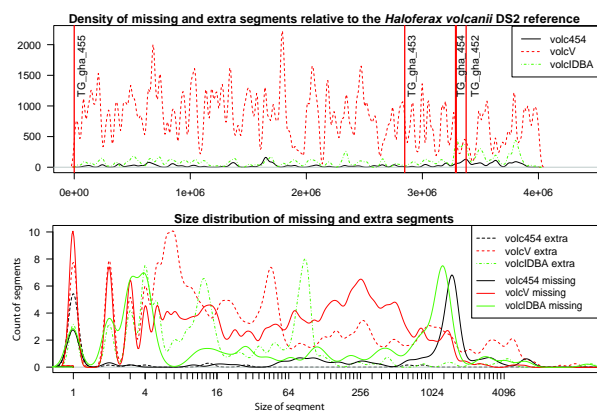


Fig. 1. Top: Density of extra and missing segments in the assemblies of *Haloferax volcanii* DS2. Reference genome coordinates are given on the x-axis, and red vertical bars delineate the boundaries of the five circular replicons in the reference genome. **Bottom:** Size distribution of missing and extra segments in each assembly. The size of a missing segment is given on the x-axis, and the count of missing segments at that size on the y-axis.




Fig. 2. Biased errors in the base calling of each assembly. Errors are not uniformly random in any of the three assemblies. See Supplementary Material for more details.

ACKNOWLEDGEMENTS

This work was supported by National Science Foundation award ER 0949453.

REFERENCES

- Altschul, S. F., Madden, T. L., Schaffer, A. A., Zhang, J., Zhang, Z., Miller, W., and Lipman, D. J. (1997). Gapped BLAST and PSI-BLAST: a new generation of protein database search programs. *Nucleic Acids Research*, **25**(17), 3389–3402.
- Bergeron, A., Mixtacki, J., and Stoye, J. (2006). A unifying view of genome rearrangements. In P. Bucher and B. M. E. Moret, editors, *WABI '06: Proceedings of the Sixth International Workshop on Algorithms in Bioinformatics*, volume 4175 of *Lecture Notes in Computer Science*, pages 163–173. Springer.
- Darling, A. C. E., Mau, B., Blattner, F. R., and Perna, N. T. (2004). Mauve: multiple alignment of conserved genomic sequence with rearrangements. *Genome Res*, **14**(7), 1394–403.
- Darling, A. E., Mau, B., and Perna, N. T. (2010). progressiveMauve: Multiple Genome Alignment with Gene Gain, Loss and Rearrangement. *PLoS ONE*, **5**(6), e11147.
- Hartman, A. L., Norais, C., Badger, J. H., Delmas, S., Haldenby, S., Madupu, R., Robinson, J., Khouri, H., Ren, Q., Lowe, T. M., Maupin-Furlow, J., Pohlschroder, M., Daniels, C., Pfeiffer, F., Allers, T., and Eisen, J. A. (2010). The Complete Genome Sequence of *Haloferax volcanii* DS2, a Model Archaeon. *PLoS ONE*, **5**(3), e9605.
- Kurtz, S., Phillippy, A., Delcher, A. L., Smoot, M., Shumway, M., Antonescu, C., and Salzberg, S. L. (2004). Versatile and open software for comparing large genomes. *Genome Biol*, **5**(2), R12.
- Phillippy, A., Schatz, M., and Pop, M. (2008). Genome assembly forensics: finding the elusive mis-assembly. *Genome Biology*, **9**, 1–13. 10.1186/gb-2008-9-3-r55.
- Rissman, A. I., Mau, B., Biehl, B. S., Darling, A. E., Glasner, J. D., and Perna, N. T. (2009). Reordering contigs of draft genomes using the Mauve Aligner. *Bioinformatics*, **25**(16), 2071–2073.